

Machine Learning: Introduction to Linear Regression, Logistic Regression, and Neural Networks

Chapter 6: Brief Introduction to Tensorflow

6.1 Tensorflow for MNIST Digits Classification

6.2 Tensorflow on a GPU Demo

Tensorflow for MNIST Digits Classification

Goal of this Section:

- Show how Tensorflow can be used for the MNIST Digits Classification problem

What is Tensorflow?

From the Tensorflow website:

- TensorFlow is an end-to-end open source platform for machine learning. It has a comprehensive, flexible ecosystem of tools, libraries and community resources that lets researchers push the state-of-the-art in ML and developers easily build and deploy ML powered applications.
- See website for details, examples, tutorials, etc <https://www.tensorflow.org/>
- Some key features:
 - Can easily set up neural networks, train, and predict
 - Version of Tensorflow optimized to make use of GPU chips to speed up training
- Many online resources available:
 - Courses on Udemy and Coursera
 - Tutorials on Youtube and various blog sites

Tensorflow for MNIST Digits Classification

- This section shows how Tensorflow can be used for the MNIST Digits Classification problem discussed in the last chapter
- See:
 - Course Version: `IntroML/Code/Version5.1/driver_neuralnetwork_mnist.py`
 - Tensorflow Version: `IntroML/Code/Version5.1/driver_tensorflow_mnist.py`
- Recall components of driver:
 1. Data loading/preparation
 2. Neural Network Definition
 3. Compilation
 4. Training
 5. Prediction
- Following slides show side-by-side comparison of course code and tensorflow versions for each component

Code Comparison: Data Loading/Preparation

- Import
 - Course code: import NeuralNetwork
 - Tensorflow: import tensorflow as tf
- Data loading/preparation:
 - Course code: load_mnist outputs feature matrices (# features x # samples) and label vectors (1 x # samples)
 - Tensorflow: requires samples axis to be along rows so take transpose using .T functionality for numpy arrays

Course Code Driver

Tensorflow Driver

```
import load_mnist
import NeuralNetwork
import matplotlib.pyplot as plt
import numpy as np
import metrics
import plot_results
import time

# (1) Set up data
ntrain = 6000
nvalid = 1000
nclass = 10
Xtrain, Ytrain, Xvalid, Yvalid = load_mnist.load_mnist(ntrain, nvalid)
```

```
import load_mnist
import matplotlib.pyplot as plt
import metrics
import numpy as np
import plot_results
import tensorflow as tf
import time

# (1) Set up data
ntrain = 6000
nvalid = 1000
nclass = 10
Xtrain, Ytrain, Xvalid, Yvalid = load_mnist.load_mnist(ntrain, nvalid)
# take transpose of inputs for tensorflow - sample axis along rows
XtrainT = Xtrain.T
YtrainT = Ytrain.T
XvalidT = Xvalid.T
YvalidT = Yvalid.T
```

Code Comparison: Neural Network Definition

- Neural Network Definition Tensorflow:
 - `tf.keras.models.Sequential` method is used to build neural network as a sequence of layers
 - `tf.keras.layers.Dense` is equivalent to the `add_layer` method from this course
 - `input_shape(784,)` in first layer defines number of features
 - 128 & `nclass` are number of units
 - `kernel_regularizer` defines regularization (can use l1 or l2) and `lamb` is coefficient
 - activation specifies activation function

Course Code Driver

```
# (2) Define model
np.random.seed(10)
lamb = 0.0
model = NeuralNetwork.NeuralNetwork(784)
model.add_layer(128,"tanh",lamb)
model.add_layer(nclass,"softmax",lamb)
# (3) Compile model
```

Tensorflow Driver

```
# (2) Define model
lamb = 0.0
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(128, input_shape=(784,), activation="tanh", kernel_regularizer=tf.keras.regularizers.l2(lamb)),
    tf.keras.layers.Dense(nclass, activation="softmax", kernel_regularizer=tf.keras.regularizers.l2(lamb))]
# (3) Compile model
```

Code Comparison: Compilation

- Neural Network Definition Tensorflow:
 - Adam optimizer specified by `tf.keras.optimizers.Adam` function
 - Loss function for multi-class classification specified as “`sparse_categorical_crossentropy`”
 - Must specify “accuracy” in metrics input to ensure accuracy is computed for each epoch

Course Code Driver

```
model.compile(loss=loss, optimizer=optimizer)
# (3) Compile model
optimizer = {"method": "Adam", "learning_rate": 0.02, "beta1": 0.9, "beta2": 0.999, "epsilon": 1e-7}
model.compile("crossentropy", optimizer)
model.summary()
```

Tensorflow Driver

```
# (3) Compile model
optimizer = tf.keras.optimizers.Adam(lr=0.02, beta_1=0.9, beta_2=0.999, epsilon=1e-7)
model.compile(optimizer=optimizer, Loss="sparse_categorical_crossentropy", metrics=["accuracy"])
model.summary()
```


Code Comparison: Compilation

- Training in Tensorflow:
 - Use fit method to perform training
 - Same inputs as in course code (training data, epochs, batch_size, validation_data)
 - Use transposed feature matrix and label vector for training dataset
 - Use transposed feature matrix and label vector for validation dataset

Course Code Driver

```
# (4) Train model
epochs = 40
time_start = time.time()
history = model.fit(Xtrain,Ytrain,epochs,batch_size=ntrain,validation_data=(Xvalid,Yvalid))
time_end = time.time()
print("Train time: {}".format(time_end - time_start))
```

Tensorflow Driver

```
# (4) Train model
epochs = 40
time_start = time.time()
history = model.fit(XtrainT,YtrainT,epochs=epochs,batch_size=ntrain,validation_data=(XvalidT,YvalidT))
time_end = time.time()
print("Train time: {}".format(time_end - time_start))
```

Code Comparison: Prediction

- Prediction in Tensorflow:
 - Use predict method to predict results for validation data set
 - predict method outputs result of activation $A^{[2]}$ (#samples x 10) at final layer
 - Use argmax to find predicted class
 - Yvalid_pred: Use expand_dims to convert to (1 x #samples) matrix
 - history.history is dictionary of results from training
 - "loss"/"acc" has history of training dataset loss/accuracy
 - "val_loss"/"val_acc" has history of validation dataset loss/accuracy

Course Code Driver

```
# (5) Predictions and plotting
# confusion matrix
Yvalid_pred = model.predict(Xvalid)
metrics.confusion_matrix(Yvalid,Yvalid_pred,nclass)
# plot loss, accuracy, and animation of results
plot_results.plot_results_history(history,["loss","valid_loss"])
plot_results.plot_results_history(history,["accuracy","valid_accuracy"])
plot_results.plot_results_mnist_animation(Xvalid,Yvalid,Yvalid_pred,100)
plt.show()
```

Tensorflow Driver

```
# (5) Predictions and plotting
Yvalid_pred = np.expand_dims(np.argmax(model.predict(XvalidT),axis=1),axis=0)
metrics.confusion_matrix(Yvalid,Yvalid_pred,nclass)
# plot loss, accuracy, and animation of results
plot_results.plot_results_history(history.history,["loss","val_loss"])
plot_results.plot_results_history(history.history,["acc","val_acc"])
plot_results.plot_results_mnist_animation(Xvalid,Yvalid,Yvalid_pred,100)
plt.show()
```

Chapter 6: Introduction to Tensorflow

6.1 Using Tensorflow for MNIST Digits Classification

6.2 Tensorflow on a GPU Demo