

# Introduction to Machine Learning: Linear and Logistic Regression and Neural Networks Using Python

# Chapter 8 Introduction to Tensorflow

# Machine Learning Frameworks

- A machine learning framework consists of building blocks for designing, training and validating deep neural networks, through a high level programming interface.
- There are many machine learning frameworks, including Tensorflow, Pytorch, Matlab, sklearn, etc. See following link for a listing:

[https://en.wikipedia.org/wiki/Comparison\\_of\\_deep-learning\\_software](https://en.wikipedia.org/wiki/Comparison_of_deep-learning_software)

- Why go through all the trouble of learning the math and algorithms and then building a basic framework from scratch in this course, when one can use one of these frameworks “off the shelf”?
- As mentioned at the beginning of the course, it is difficult to truly understand what a framework is doing without going through the math and algorithms and doing the coding.

# What is Tensorflow?

From the Tensorflow website:

- TensorFlow is an end-to-end open source platform for machine learning. It has a comprehensive, flexible ecosystem of tools, libraries and community resources that lets researchers push the state-of-the-art in ML and developers easily build and deploy ML powered applications.
- See website for details, examples, tutorials, etc <https://www.tensorflow.org/>
- Some key features:
  - Can easily set up neural networks, train, and predict
  - There is a version of Tensorflow optimized to make use of GPU chips, that are faster than CPU for training
- Many online resources available:
  - Courses on Udemy and Coursera
  - Tutorials on Youtube and various blog sites

# Introduction to Tensorflow

Section	Title	Description
8.1	Tensorflow for MNIST Digits Classification	This section shows how to use Tensorflow for the MNIST digits classification problem. We will create a driver analogous to driver_casestudy_mnist.
8.2	Demo of Tensorflow on a GPU	This section present a demo of running Tensorflow on a GPU for the MNIST digits classification problem

# 8.1 Tensorflow for MNIST Digits Classification

# Tensorflow for MNIST Digits Classification

Goal of this Section:

- Show how Tensorflow can be used for MNIST Digits Classification

# Tensorflow for MNIST Digits Classification

Recall components of driver:

1. Data loading/preparation
  2. Neural Network Definition
  3. Compilation
  4. Training
  5. Prediction
- Following slides show side-by-side comparison of course code and tensorflow drivers
    - Course Driver: `driver_casestudy_mnist.py`
    - Tensorflow Driver: `driver_casestudy_mnist_tensorflow.py`



# Import Packages

Course Code Driver

```
import load_mnist
import NeuralNetwork
import matplotlib.pyplot as plt
import numpy as np
import metrics
import Optimizer
import plot_results
import time
```

Tensorflow Driver

```
import load_mnist
import tensorflow as tf
import matplotlib.pyplot as plt
import numpy as np
import metrics
import onehot
import Optimizer
import plot_results
import time
```

# Code Comparison: Data Loading/Preparation

- Data loading/preparation:
  - Course Code: load\_mnist outputs feature matrices (# features x # samples) and label vectors (1 x # samples)
  - Tensorflow: requires samples axis to be along rows (feature matrix should be #samples x # features) so take transpose using .T functionality for numpy arrays

Course Code Driver

```
# (1) Set up data
ntrain = 60000
nvalid = 10000
nclass = 10
Xtrain,Ytrain,Xvalid,Yvalid = load_mnist.load_mnist(ntrain,nvalid)
```

Tensorflow Driver

```
# (1) Set up data
ntrain = 60000
nvalid = 10000
nclass = 10
Xtrain,Ytrain,Xvalid,Yvalid = load_mnist.load_mnist(ntrain,nvalid)
# Take transpose of inputs for tensorflow - sample axis along rows
XtrainT = Xtrain.T
YtrainT = Ytrain.T
XvalidT = Xvalid.T
YvalidT = Yvalid.T
```

# Code Comparison: Neural Network Definition

- Neural Network Definition Tensorflow:
  - `tf.keras.models.Sequential` method is used to build neural network as a sequence of layers
  - `tf.keras.layers.Dense` is equivalent to the `add_layer` method from this course
  - `input_shape(nfeature, )` in first layer defines number of features
  - 128 & `nclass` are number of units
  - `kernel_regularizer` defines regularization (can use l1 or l2) and `lamb` is coefficient
  - `activation` specifies activation function

## Course Code Driver

```
# (2) Define model
nfeature = Xtrain.shape[0]
np.random.seed(10)
lamb = 0.0
model = NeuralNetwork.NeuralNetwork(nfeature)
model.add_layer(128,"relu",lamb)
model.add_layer(nclass,"softmax",lamb)
```

## Tensorflow Driver

```
# (2) Define model
nfeature = Xtrain.shape[0]
lamb = 0.0
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(nfeature, activation="relu", kernel_regularizer=tf.keras.regularizers.l2(lamb)),
    tf.keras.layers.Dense(nclass, activation="softmax", kernel_regularizer=tf.keras.regularizers.l2(lamb))])
```

# Code Comparison: Compilation

- Neural Network Definition Tensorflow:
  - Adam optimizer specified by `tf.keras.optimizers.Adam` function
  - Loss function for multi-class classification specified as `"sparse_categorical_crossentropy"`
  - Must specify `"accuracy"` in metrics input to ensure accuracy is computed for each epoch
  - `summary()` produces summary of layers and number of parameters

## Course Code Driver

```
# (3) Compile model
optimizer = Optimizer.Adam(0.02,0.9,0.999,1e-7)
model.compile("crossentropy",optimizer)
model.summary()
# (4) Train model
```

## Tensorflow Driver

```
# (3) Compile model
optimizer = tf.keras.optimizers.Adam(lr=0.02, beta_1=0.9, beta_2=0.999, epsilon=1e-7)
model.compile(optimizer=optimizer, Loss="sparse_categorical_crossentropy", metrics=["accuracy"])
model.summary()
```

# Code Comparison: Training

- Training in Tensorflow:
  - Use fit method to perform training
  - Same inputs as in course code (training data, epochs, batch\_size, validation\_data)
  - Use transposed feature matrix and label vector for training dataset
  - Use transposed feature matrix and label vector for validation dataset

## Course Code Driver

```
# (4) Train model
epochs = 40
time_start = time.time()
history = model.fit(Xtrain,Ytrain,epochs,batch_size=ntrain,validation_data=(Xvalid,Yvalid))
time_end = time.time()
print("Train time: {}".format(time_end - time_start))
```

## Tensorflow Driver

```
# (4) Train model
epochs = 40
time_start = time.time()
history = model.fit(XtrainT,YtrainT,epochs=epochs,batch_size=ntrain,validation_data=(XvalidT,YvalidT))
time_end = time.time()
print("Train time: {}".format(time_end - time_start))
```

# Code Comparison: Prediction

- Prediction in Tensorflow:
  - Use predict method to predict results for validation data set
    - predict method outputs result of activation  $A^{[2]}$  (#samples x 10) at final layer
    - Take transpose and use onehot\_inverse to compute class for each sample
  - history.history is dictionary of results from training
    - “loss”/”accuracy” has history of training dataset loss/accuracy
    - “val\_loss”/”val\_accuracy” has history of validation dataset loss/accuracy

## Course Code Driver

```
# (5) Predictions and plotting
Yvalid_pred = model.predict(Xvalid)
metrics.confusion_matrix(Yvalid,Yvalid_pred,nclass)
# plot loss, accuracy, and animation of results
plot_results.plot_results_history(history,["loss","valid_loss"])
plot_results.plot_results_history(history,["accuracy","valid_accuracy"])
plot_results.plot_results_mnist_animation(Xvalid,Yvalid,Yvalid_pred,
    model.get_Afinal(),100)
plt.show()
```

## Tensorflow Driver

```
# (5) Predictions and plotting
Afinal = model.predict(XvalidT).T
Yvalid_pred = onehot.onehot_inverse(Afinal)
metrics.confusion_matrix(Yvalid,Yvalid_pred,nclass)
# plot loss, accuracy, and animation of results
plot_results.plot_results_history(history.history,["loss","val_loss"])
plot_results.plot_results_history(history.history,["accuracy","val_accuracy"])
plot_results.plot_results_mnist_animation(Xvalid,Yvalid,Yvalid_pred,Afinal,100)
plt.show()
```

# Timings for MNIST Digits Classification

Run	Machine	Chip	Code	Training Time 60K Train samples 10K Valid samples	Accuracy Train	Accuracy Valid
1	My Windows	CPU: Intel i5-6200 2.30GHz	Course Framework	65.2 seconds	91.5%	91.6%
2	My Windows	CPU: Intel i5-6200 2.30GHz	Tensorflow2.1	24.3 seconds	94.9%	94.6%

# New Code for MNIST Using Tensorflow

Function/component	Input	Description
driver_casestudy_ mnist_tensorflow		Driver



# 8.1 Tensorflow for MNIST Classification DEMO

- Code located at: IntroML/Code/Version5.1
- Course Driver: driver\_casestudy\_mnist.py
- Tensorflow Driver: driver\_casestudy\_mnist\_tensorflow.py

Course Resources at:

- <https://github.com/satishchandrareddy/IntroML/>

## 8.2 Demo of Tensorflow on a GPU

# Demo of Tensorflow on a GPU

Goal of this Section:

- Perform a Demo of Tensorflow on a GPU

# What is a GPU and Why Use for Machine Learning?

- A GPU (Graphical Processing Unit) is a specialized chip designed for image processing (used for gaming)
- Parallel structure of GPU make them more efficient than CPU for algorithms that process data in parallel
- Source:  
[https://en.wikipedia.org/wiki/Graphics\\_processing\\_unit](https://en.wikipedia.org/wiki/Graphics_processing_unit)
- Machine learning frameworks have been developed to take advantage of the features of GPU and run much faster than on a CPU

# NVIDIA GPU and Machine Learning Frameworks

- NVIDIA (GPU chip company) provides GPU support for number of machine learning frameworks including: Tensorflow, Pytorch, Matlab, etc. See

<https://developer.nvidia.com/deep-learning-frameworks>

- See following for instructions on installing GPU version of Tensorflow

<https://www.tensorflow.org/install/gpu>

# Tensorflow on a GPU

- In previous section, we showed a Tensorflow code for MNIST classification
- No change to code is required to run on GPU
- In Demo will run Tensorflow on Azure VM with/without GPU
  - Tensorflow will automatically run with GPU turned on
  - Add following before import of tensorflow to turn off GPU and use CPU  
import os  
os.environ["CUDA\_VISIBLE\_DEVICES"] = "-1"

# Results MNIST Digits Classification

Run	Machine	Chip	Code	Training Time 40 Epochs 60K Train samples 10K Valid samples	Accuracy Train	Accuracy Valid
1	My Windows	CPU: Intel i5-6200 2.30GHz	Course Framework	65.2 seconds	91.2%	90.6%
2	My Windows	CPU: Intel i5-6200 2.30GHz	Tensorflow2.1	24.3 seconds	93.2%	92.9%
3	Azure VM	CPU: Intel Xeon E5-2690 v3 2.60GHz	Course Framework	45.0 seconds	91.2%	90.6%
4	Azure VM	CPU: Intel Xeon E5-2690 v3 2.60GHz	Tensorflow2.1	13.5 seconds	93.2%	92.7%
5	Azure VM	GPU: NVIDIA Tesla M60	Tensorflow2.1	7.9 seconds	93.5%	93.3%